

Assignment Project Exam Help

Machine learning lecture slides

COMS 4771 Fall 2020

<https://powcoder.com>

Add WeChat powcoder

Nearest neighbor classification

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- ▶ Optical character recognition (OCR) example
- ▶ Nearest neighbor rule
- ▶ Error rate, test error rate
- ▶ k -nearest neighbor rule
- ▶ Hyperparameter tuning via cross-validation
- ▶ Distance functions, features
- ▶ Computational issues

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Example: OCR for digits


- ▶ Goal: Automatically label images of handwritten digits
- ▶ Possible labels are $\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$
- ▶ Start with a large collection of already-labeled images

- ▶ $D := \{(x_1, y_1), \dots, (x_n, y_n)\}$

- ▶ x_i is the i -th image; $y_i \in \{0, 1, \dots, 9\}$ is the corresponding label.

- ▶ The National Institute for Standards and Technology (NIST) has amassed such a data set.

Add WeChat powcoder



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Figure 1: Some images of handwritten digits from MNIST data set

Nearest neighbor (NN) classifier

- ▶ Nearest neighbor (NN) classifier NN_D :

- ▶ Represented using collection of labeled examples

$D := ((x_1, y_1), \dots, (x_n, y_n))$, plus a snippet of code

- ▶ Input: x

- ▶ Find x_i in D that is “closest” to x (the nearest neighbor)

- ▶ (Break ties in some arbitrary fixed way)

- ▶ Return: y_i

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Naïve distance between images of handwritten digits (1)

- ▶ Treat (grayscale) images as vectors in Euclidean space \mathbb{R}^d
 - ▶ $d = 28^2 = 784$
 - ▶ Generalizes physical 3-dimensional space
- ▶ Each point $x = (x_1, \dots, x_d) \in \mathbb{R}^d$ is a vector of d real numbers
 - ▶ $\|x - z\|_2 = \sqrt{\sum_{j=1}^d (x_j - z_j)^2}$
 - ▶ Also called ℓ_2 distance
 - ▶ WARNING: Here, x_j refers to the j -th coordinate of x

Add WeChat powcoder




Figure 2: Grayscale pixel representation of an image of a handwritten “4”

Naïve distance between images of handwritten digits (2)

- ▶ Why use this for images?

Assignment Project Exam Help

<https://powcoder.com>

- ▶ Why not use this for images?

Add WeChat powcoder

- ▶ What is the core prediction problem?

Assignment Project Exam Help

- ▶ What features (i.e., predictive variables) are available?

- ▶ Will these features be available at time of prediction?

<https://powcoder.com>

- ▶ Is there enough information (“training data”) to learn the relationship between the features and label?

Add WeChat powcoder

- ▶ What are the modeling assumptions?
- ▶ Is high-accuracy prediction a useful goal for the application?

- ▶ Error rate (on a collection of labeled examples S)
 - ▶ Fraction of labeled examples in S that have incorrect label prediction from \hat{f}
 - ▶ Written as $\text{err}(\hat{f}, S)$
 - ▶ (Often, the word “rate” is omitted)
- ▶ OCR via NN:

<https://powcoder.com>

Add WeChat powcoder

Test error rate (1)

- ▶ Better evaluation: test error rate

- ▶ Train/test split, $S \cap T = \emptyset$

- ▶ S is training data, T is test data

- ▶ Classifier \hat{f} only based on S

- ▶ Training error rate: $\text{err}(\hat{f}, S)$

- ▶ Test error rate: $\text{err}(\hat{f}, T)$

- ▶ On OCR data, test error rate is $\text{err}(\text{NN}_S, T) = 3.09\%$

- ▶ Is this good?

- ▶ What is the test error rate of uniformly random predictions?

- ▶ What is the test error rate of a constant prediction?


- ▶ Why is test error rate meaningful?

Assignment Project Exam Help

<https://powcoder.com>

- ▶ What are the drawbacks of evaluation via test error rate?

Add WeChat powcoder



Assignment Project Exam Help

<https://powcoder.com>

Figure 3: A test example and its nearest neighbor in training data (2, 8)

Add WeChat powcoder




Assignment Project Exam Help

<https://powcoder.com>

Figure 4: A test example and its nearest neighbor in training data (3, 5)

Add WeChat powcoder



Assignment Project Exam Help

<https://powcoder.com>

Figure 5: A test example and its nearest neighbor in training data (5, 4)

Add WeChat powcoder



Assignment Project Exam Help

<https://powcoder.com>

Figure 6: A test example and its nearest neighbor in training data (4, 1)

Add WeChat powcoder

More on the modeling assumptions

- ▶ Modeling assumption: Nearby images are more likely to have the same label than different labels.

Assignment Project Exam Help

- ▶ This is an assumption about the choice of distance function.
- ▶ In our OCR example, this is an assumption about the choice of features

<https://powcoder.com>

Add WeChat powcoder

- ▶ What are the kinds of errors made by NN_S ?

Assignment Project Exam Help

<https://powcoder.com>

Figure 7: A test example and its nearest neighbor in training data (2, 8)

Add WeChat powcoder

Figure 8: Three nearest neighbors of the test example (8,2,2)

Upgrade: k -NN

- ▶ k -nearest neighbor (k -NN) classifier $\text{NN}_{k,D}$

- ▶ Input: x

- ▶ Find the k nearest neighbors of x in L
- ▶ Return the plurality of the corresponding labels

- ▶ As before, break ties in some arbitrary fixed way

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Typical effect of k

- ▶ Smaller k : smaller training error rate
- ▶ Larger k : higher training error rate, but predictions more “stable” due to voting
- ▶ On OCR data: lowest test error rate achieved at $k = 3$

k	1	3	5	7	9
$\text{err}(\text{NN}_k, S; T)$	0.0309	0.0295	0.0312	0.0306	0.0341

Add WeChat powcoder

Hyperparameter tuning

- ▶ k is a hyperparameter of k -NN
- ▶ How to choose hyperparameters?
 - ▶ Bad idea: Choosing k that yields lowest training error rate (degenerate choice: $k = 1$)
 - ▶ Better idea: Simulate train/test split on the training data
- ▶ Hold-out validation
 - ▶ Randomly split D into A and B
 - ▶ Compute validation error rate for all $k \in \{1, 3, 5, 7, 9\}$:

$$V_k := \text{err}(\text{NN}_{k,A}, B)$$

- ▶ Let \hat{k} be the value of k for which V_k is smallest
- ▶ Classifier to use is $\text{NN}_{\hat{k},S}$

Upgrade: Distance functions (1)

- ▶ Specialize to input types
 - ▶ Edit distance for strings
 - ▶ Shape distance for images
 - ▶ Time warping distance for audio waveforms

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Upgrade: Distance functions (2)

- ▶ Generic distances for vectors of real numbers

- ▶ ℓ_p distances

Assignment Project Exam Help

$$\|x - z\|_p = \left(\sum_{j=1}^d |x_j - z_j|^p \right)^{1/p}.$$

<https://powcoder.com>

- ▶ What are the unit balls for these distances (in \mathbb{R}^d)?

Add WeChat powcoder

Upgrade: Distance functions (3)

- Distance functions for images of handwritten digits

	distance	l_2	tangent	shape
test error rate	0.0309	0.0283	0.0110	0.0065

<https://powcoder.com>

Add WeChat powcoder

Features

- ▶ When using numerical features (arranged in a vector from \mathbb{R}^d):
 - ▶ Scale of features matters
 - ▶ Noisy features can ruin NN
 - ▶ E.g., consider what happens in OCR example if you have another 10000 additional features that are pure “noise”
 - ▶ Or a single pure noise feature whose scale is $10000\times$ the nominal scale of pixel values

Assignment Project Exam Help

<https://powcoder.com>

- ▶ Add WeChat powcoder
- ▶ “Curse of dimension”
 - ▶ Weird effects in \mathbb{R}^d for large d
 - ▶ Can find $2^{\Omega(d)}$ points that are approximately equidistant

- ▶ Brute force search: $\Theta(dn)$ time for each prediction (using Euclidean distance in \mathbb{R}^d)
- ▶ Clever data structures: “improve” to $2^d \lg(n)$ time
- ▶ Approximate nearest neighbors: sub-linear time to get “approximate” answers
- ▶ E.g., find point among the top-1% of closest points?

<https://powcoder.com>

Add WeChat powcoder